

Live E2E Report Check — Analytics Report

Business Analytics Report

Project: Live E2E Report Check

Dataset: sample_customer_churn.csv

Analysis date: 2026-09-24

Prepared by: Report Verify

Generated by the AI Data Scientist Platform

Table of Contents

1. Executive Summary	4
Key verified findings	4
Top recommendations	4
2. About the Dataset	5
Column dictionary	5
3. Data Quality & Cleaning	6
Original dataset quality (as uploaded)	6
Outliers	6
Cleaning performed	6
4. Key EDA Findings	7
Trend over time (Signup Date)	7
Churn breakdown	7
5. Modeling Approach	8
Features used	8
Features excluded, and why	8
6. Model Performance	9
Model comparison	9
In plain language	9
Confusion Matrix	9
7. What Drives the Outcome	10
8. Future Outlook	11
9. AI Business Insights AI-generated	12
Key Business Insights	12
Recommendations	12
Areas to Investigate	12
Potential Risks	12
Suggested Next Actions	12
10. Limitations	13
11. Appendix	14
Full model metrics	14
Preprocessing	14
Hyperparameters	14
Dataset statistics	14

1. Executive Summary

This report analyzes **sample_customer_churn.csv** (300 rows, 12 columns) covering 2024-01-01 to 2024-12-28 to answer the business question: *final live verification*.

Key verified findings

- The original dataset contained 0 missing value(s) and 0 duplicate row(s), which have since been cleaned.
- A Random Forest model was trained to predict **Churn** (classification). The model correctly classifies about 71.7% of cases (Accuracy).

Model reliability, in plain terms: The model correctly classifies about 71.7% of cases (Accuracy). F1 score (balancing precision and recall) is 0.713. ROC-AUC is 0.792 (1.0 = perfect separation, 0.5 = no better than chance).

Top recommendations

- The most common observed outcome for Churn is '1' (32 of 60 test cases).
- Among customers, the 'Tenure Months: Lower' group shows the highest observed rate (77.0%, n=100), compared to 'Tenure Months: Higher' (33.0%, n=100).

2. About the Dataset

`sample_customer_churn.csv` contains 300 rows and 12 columns. It spans the period 2024-01-01 to 2024-12-28 (based on Signup Date).

Business question: final live verification

Column dictionary

Column	Type	Missing	Notes
Customer Id	Identifier	0.0%	
Age	Numeric	0.0%	typical range 18.00–75.00
Region	Category	0.0%	most common: West
Plan	Category	0.0%	most common: Standard
Contract Type	Category	0.0%	most common: Month-to-month
Tenure Months	Numeric	0.0%	typical range 1.00–72.00
Monthly Charge	Numeric	0.0%	typical range 15.27–118.69
Total Spend	Numeric	0.0%	typical range 75.45–8,521.1
Support Tickets	Numeric	0.0%	typical range 0.00–10.00
Satisfaction Score	Numeric	0.0%	typical range 1.00–4.90
Signup Date	Date/Time	0.0%	
Churn	Numeric	0.0%	typical range 0.00–1.00

3. Data Quality & Cleaning

Original dataset quality (as uploaded)

Total rows	Missing cells	Duplicate rows	Issues detected
300	0	0	1

Outliers

Column	Finding	Affected
total_spend	'total_spend' has 8 statistical outlier(s) outside [-2921.19, 7084.38].	8



Cleaning performed

The dataset was marked cleaned, but no specific actions were logged.

4. Key EDA Findings

Total Rows	Churn Rate	Total Age	Total Tenure Months
300	53.7%	14,758	11,137

Trend over time (Signup Date)



Churn breakdown

Average: 0.54 (range 0.0–1.0)

Value	Count	Percent
1	161	53.7%
0	139	46.3%

5. Modeling Approach

A **classification** model was trained to predict **Churn**, using 11 feature(s).

Features used

Age, Region, Plan, Contract Type, Tenure Months, Monthly Charge, Total Spend, Support Tickets, Signup Date: Year, Signup Date: Month, Signup Date: Dayofweek

Features excluded, and why

Column	Reason excluded
Customer Id	Identifier column (e.g. an ID) — carries no predictive signal.
Signup Date	Raw date — replaced with derived year / month / day-of-week features instead.
Satisfaction Score	Excluded by default 'satisfaction_score' often represents information only available AFTER the outcome occurs (e.g. a rating or tip given once an order is complete), and shows a real (23% correlated) relationship with the target here. Review whether this would actually be known at prediction time — if not, exclude it below.

Technical details

Algorithm: Random Forest · Preprocessing: 8 numeric + 3 categorical feature(s), scaled/encoded · Train/test split: 20.0% held out · Cross-validation folds: 5 · Random seed: 42

6. Model Performance

Model comparison


Model	Accuracy	Precision	Recall	F1 Score	ROC-AUC
Random Forest	0.7167	0.7197	0.7167	0.7130	0.7919

In plain language


- The model correctly classifies about 71.7% of cases (Accuracy).
- F1 score (balancing precision and recall) is 0.713.
- ROC-AUC is 0.792 (1.0 = perfect separation, 0.5 = no better than chance).

Compared against a trivial baseline (F1 Score = 0.371): this model clearly outperforms it.

Confusion Matrix



7. What Drives the Outcome



Feature	Relative importance
Tenure Months	18.9%
Total Spend	14.2%
Support Tickets	12.2%
Monthly Charge	10.1%
Age	10.1%
Contract Type: Month to month	7.8%
Signup Date: Month	5.6%
Contract Type: Two year	4.8%
Signup Date: Dayofweek	4.7%
Contract Type: One year	1.8%

8. Future Outlook

- The most common observed outcome for Churn is '1' (32 of 60 test cases).
- Among customers, the 'Tenure Months: Lower' group shows the highest observed rate (77.0%, n=100), compared to 'Tenure Months: Higher' (33.0%, n=100).

Group	Observed rate	Group size
Tenure Months: Lower	77.0%	100
Tenure Months: Middle	51.0%	100
Tenure Months: Higher	33.0%	100

9. AI Business Insights AI-generated

Key Business Insights

- Customer tenure is the primary driver of retention, with lower-tenure customers exhibiting a 77.0% churn rate compared to 33.0% for higher-tenure customers.
- Total spend, support tickets, and monthly charges heavily influence churn risks, representing top predictive factors alongside customer age.
- The underlying dataset was clean, containing 0 missing cells and 0 duplicate rows both originally and currently across 300 records, supporting a Random Forest model with 71.67% accuracy and a 0.7919 ROC AUC.

Recommendations

- A possible response is to enhance the early-stage onboarding experience to help stabilize retention during the critical initial months.
- Implementing proactive outreach for clients generating support tickets may reduce dissatisfaction before accounts decide to churn.
- Reviewing pricing models or offering targeted incentives for accounts with high monthly charges might assist in retaining price-sensitive users.

Areas to Investigate

- Investigating specific support ticket categories may reveal recurring product pain points or service bottlenecks.
- Analyzing the exact transition period where churn drops from 77.0% to 33.0% could help pinpoint key retention milestones.

Potential Risks

- High early-tenure customer loss presents a risk to overall customer lifetime value and marketing acquisition efficiency.
- Dissatisfaction linked to support ticket volume or elevated monthly charges may increase susceptibility to competitor offerings.

Suggested Next Actions

- Segment support ticket data to identify common issues affecting newer accounts.
- Test retention workflows targeted at early-tenure accounts to improve long-term customer stability.

Generated by gemini. Grounded strictly in the verified results above — the AI does not calculate or invent any numbers.

10. Limitations

- Cross-validation scores varied by ± 0.042 across folds — treat the headline metric as an estimate with this much natural uncertainty, not an exact figure.
- Model predictions are estimates based on historical patterns and may not hold if underlying business conditions change. They should inform, not replace, domain expertise.

11. Appendix

Full model metrics

Metric	Value
Accuracy	0.7167
Precision	0.7197
Recall	0.7167
F1 Score	0.713
ROC-AUC	0.7919
Cv Mean	0.7648
Is Weak	False

Preprocessing

Setting	Value
Test Size	0.2
Cv Folds	5
Random seed	42
Model version	1

Hyperparameters

Parameter	Value
Random Seed	42

Dataset statistics

Statistic	Value
Rows	300
Columns	12